

Paper Coding Instructions

1 Artifact Selection Criteria

Papers are selected based on adoption by the following artifacts:

1.1 Tools (Criterion: $\geq 1,000$ GitHub stars)

Tool	Stars	Description
CleverHans	6,401	Adversarial example library (Toronto)
IBM ART	5,789	Adversarial Robustness Toolbox
TextAttack	3,348	NLP adversarial attacks
PyRIT	3,343	LLM red-teaming (Microsoft)
Foolbox	2,936	Adversarial attacks (Bethge Lab)

1.2 Benchmarks (Criterion: Peer-reviewed publication)

Benchmark	Venue	Description
RobustBench	NeurIPS 2021	Adversarial robustness leaderboard
AutoAttack	ICML 2020	Standardized attack evaluation
HarmBench	ICML 2024	LLM jailbreak evaluation

1.3 Regulatory (Criterion: Industry threat framework)

Framework	Description
MITRE ATLAS	Adversarial ML tactics/techniques (like ATT&CK)

2 Paper Selection Criteria

71 papers selected for coding:

1. **61** papers cited by 2+ artifacts (strongest adoption evidence)
2. **10** papers from MITRE ATLAS only (regulatory adoption)

3 The 12 Coding Columns

3.1 Group 1: Basic Info (G1–G7)

Col	Question	Options	How to Decide
G1	Is this an attack, defense, or evaluation?	Attack / Defense / Evaluation	Read abstract. What did they build?
G2	What type of attack?	Evasion / Poisoning / Privacy / N/A	Evasion = fool model at test time. Poisoning = corrupt training data. Privacy = steal data/model. N/A = defense papers
G3	What domain?	Vision / NLP / Malware / Audio / Tabular / LLM / Cross-domain	What data did they test on? ImageNet = Vision. Text = NLP. ChatGPT = LLM.
G4	Where published?	ML / Security / Journal / arXiv-only	ML = NeurIPS, ICML, ICLR, CVPR, ACL. Security = S&P, CCS, USENIX, NDSS.
G5	Is code available NOW?	Yes / No	Google “paper name github”. Is there code?
G6	When was code released?	At-pub / Post-pub / Never	At-pub = within 1 month of paper. Post-pub = later.
G7	Publication year	2014–2025	Year of first public version

3.2 Group 2: Threat Model (T1–T2) — Attack Papers Only

Col	Question	Options	How to Decide
T1	How much model access?	White / Gray / Black	White = has weights/gradients. Gray = surrogate model. Black = queries only
T2	Uses gradients?	Yes / No	If they compute ∇L anywhere, it's Yes

Leave T1 and T2 blank for defense papers.

3.3 Group 3: Practical Evaluation (Q1–Q3)

Col	Question	Options	How to Decide
Q1	Tested on real system?	Yes / Partial / No	Yes = Google API, Tesla, ChatGPT. Partial = realistic sim. No = CIFAR/ImageNet only
Q2	Reported cost?	Yes / No	Did they say “X queries” or “Y seconds”?
Q3	Tested against defenses?	Yes / No / N/A	Yes = tested vs. adversarial training. N/A = for defense papers

4 Pre-filled Columns (Auto-extracted)

The following columns are already filled:

Column	Description
selection_reason	Why paper was selected (2+ artifacts / MITRE ATLAS)
is_aml_paper	Always YES for this dataset
arxiv_id	arXiv identifier
found_in_artifacts	Which artifacts cite this paper
num_artifacts	How many artifacts cite it
paper_title	Full paper title
paper_authors	Author names
paper_pub_date	Publication date (YYYY-MM-DD)
first_adoption_date	When first artifact added it
first_adoption_artifact	Which artifact adopted first
all_adoptions	All adoption events (artifact:date pairs)
adoption_lag_months	Months from paper to first adoption

5 Your Workflow

1. Open `extraction_runs/run1/papers_for_coding_71.csv`
2. For each paper, fill in: G1, G2, G3, G4, G5, G6, G7, T1, T2, Q1, Q2, Q3
3. Use arXiv link to read the paper: [https://arxiv.org/abs/\[arxiv_id\]](https://arxiv.org/abs/[arxiv_id])
4. Estimated time: 20–30 minutes per paper

6 Files

All files are in `extraction_runs/run1/`:

File	Purpose
<code>papers_for_coding_71.csv</code>	71 papers for manual coding
<code>papers_all.csv</code>	Full 277 AML papers (for statistics)
<code>papers_by_artifact.csv</code>	Sorted by artifact (for verification)